

SpaceX Falcon 9 First-Stage Landing Prediction

Evidence and interpretation

Completed presentation in PDF format. This report documents the complete IBM Data Science Capstone workflow and provides substantive methods, results, visuals, interpretation, and conclusions.

Data Science Capstone Project Report

SpaceX Falcon 9 First-Stage Landing Prediction

Completed presentation slides in PDF format

IBM Data Science Capstone | Dataset period: 2010-2020

Supporting visual from the completed analysis

GitHub Repository and Complete Presentation Evidence

Q1 / Rubric 1.1 - YES. Public GitHub repository: <https://github.com/TristinWen/spacex-falcon9-landing-prediction>. It contains all seven completed notebooks plus dashboard/spacex-dash-app.py, dashboard CSV/screenshots, the Folium HTML map, and this final PDF.

Completed notebooks: API collection; Wikipedia web scraping; data wrangling; EDA visualization; SQLite SQL analysis; Folium launch-site analysis; and machine-learning classification.

Q2 / Rubric 1.2 - YES. This document is the completed 27-page presentation submitted in PDF format; every following page contains native text plus supporting analysis visuals.

Q3 / Rubric 1.3 - YES. Executive Summary on page 3 explains all methods and reports 90 launches, 41.7% KSC share, 83.3% test accuracy, and 92.9% best cross-validation.

Q4 / Rubric 1.4 - YES. Introduction on page 4 explains reusable Falcon 9 economics, the research problem, predictive objective, and binary Class 1/Class 0 target.

Q5 / Rubrics 1.5-1.7 - YES. Pages 5-7 substantively explain SpaceX API collection, BeautifulSoup web scraping, missing-value treatment, feature retention, and outcome-label construction.

Q6 / Rubric 1.8 - YES. Page 8 explains why scatter, bar, and line charts were used to study payload, flight maturity, site, orbit, outcome, and yearly trends.

Q7 - YES. Predictive classification methodology on page 24 explains train/test preparation, GridSearchCV tuning, and comparison of logistic regression, SVM, decision tree, and KNN.

Q8 / Rubric 1.11 - YES. Pages 11-16 show and interpret five required scatter/bar results plus the yearly success trend; conclusions include later-flight improvement and orbit-specific effects.

Q9 / Rubrics 1.9 and 1.12 - YES. Pages 9 and 17-19 explain SQL queries and show results for sites, date ranges, payload totals/averages, first landing, outcomes, rankings, and time analysis.

Q10 / Rubric 1.13 - YES. Pages 10 and 20-21 show Folium site markers, clustered launch records, success/failure markers, and proximity to coast, city, highway, and railway.

Q11 / Rubric 1.14 - YES. Pages 10 and 22-23 show the Plotly Dash dropdown/slider, all-site success pie chart, and site-filtered payload-versus-launch-outcome scatter plot.

Q12 / Rubric 1.15 - YES. Pages 24-25 report four tuned models, cross-validation and test accuracy, and four confusion matrices with true/false prediction outcomes.

Q13 - YES. Page 26 concludes that the decision tree is the best cross-validation model and recommends drift monitoring, new variables, more data, and validation on newer launches.

Q14 - YES. Creativity is demonstrated through a consistent evidence-led visual system, interactive Dash controls, geographic proximity analysis, comparative model scorecards, and rubric mapping.

Q15 - YES. Innovative insight on page 26: payload, orbit, site, and flight maturity interact rather than acting independently; the model should support engineering judgment and be monitored for drift.

Executive Summary

Evidence and interpretation

The project collected Falcon 9 records through the SpaceX API and Wikipedia scraping, cleaned and explored the data with visualization and SQL, built Folium and Dash analytics, and compared four classification models. The prepared dataset contains 90 launches. KSC LC-39A contributes 41.7% of successful launches. All models reached 83.3% test accuracy, while the decision tree led cross-validation at 92.9%.

RUBRIC 1.3 - METHODS AND KEY RESULTS

Executive Summary

Methods

SpaceX API, Wikipedia web scraping, data wrangling, EDA visualization, SQL, Folium, Plotly Dash, and four classification models.

90

Falcon 9 launches

87.8%

prepared-label success share

Key results

KSC LC-39A contributes 41.7% of successful launches. All four classifiers achieve 83.3% test accuracy; decision tree leads cross-validation at 92.9%.

Supporting visual from the completed analysis

Introduction: Project Background and Problem

Evidence and interpretation

SpaceX can recover and reuse Falcon 9 first stages, so landing success affects launch cost and mission economics. The analysis asks which mission characteristics are associated with successful landings and whether those characteristics can predict the binary landing outcome: Class 1 for success and Class 0 for failure.

RUBRIC 1.4 - INTRODUCTION

Introduction: Project Background and Problem

Background

Falcon 9 first stages can be recovered and reused. Recovery can lower launch cost and affect competitive bid estimates.

Problem statement

Identify the mission characteristics associated with successful landing and predict whether the first stage will land successfully.

Target: Class = 1 for successful landing; Class = 0 for unsuccessful landing.

Supporting visual from the completed analysis

Data Collection - SpaceX API

Evidence and interpretation


The API workflow requested past launches, normalized JSON responses, resolved rocket, launchpad, payload, and core identifiers, assembled a pandas DataFrame, and retained Falcon 9 missions. The completed notebook produces 90 launch records and 17 analytical fields.

Repository evidence

<https://github.com/TristinWen/spacex-falcon9-landing-prediction>

RUBRIC 1.5 - API CALLS AND KEY PHRASES

Data Collection - SpaceX API



Output: 90 Falcon 9 launches and 17 analytical fields.

Key fields: payload mass, orbit, launch site, landing outcome, reuse history, grid fins, legs, block, coordinates.

GitHub URL: <https://github.com/TristinWen/spacex-falcon9-landing-prediction>

Supporting visual from the completed analysis

Data Collection - Web Scraping

Evidence and interpretation


The scraping workflow requested the historical Falcon 9 launch page, parsed the HTML with BeautifulSoup, located the launch table, extracted headers and rows, cleaned cell values, and converted the result into a pandas DataFrame containing 121 historical launch rows and 11 fields.

Repository evidence

<https://github.com/TristinWen/spacex-falcon9-landing-prediction>

RUBRIC 1.6 - SCRAPING PROCESS AND KEY PHRASES

Data Collection - Web Scraping



Output: 121 historical launch rows with 11 fields.

Fields include flight number, date, time, booster, site, payload, mass, orbit, customer, launch outcome, and landing outcome.

GitHub URL: <https://github.com/TristinWen/spacex-falcon9-landing-prediction>

Supporting visual from the completed analysis

Data Wrangling Methodology

Evidence and interpretation

Cleaning retained the required columns, replaced missing PayloadMass values with the sample mean, preserved missing landing-pad cases, counted launches by site and orbit, mapped landing outcomes to the Class label, and produced 90 records with no missing PayloadMass values and a prepared success share of 87.8%.

Repository evidence

<https://github.com/TristinWen/spacex-falcon9-landing-prediction>

RUBRIC 1.7 - DATA CLEANING AND PROCESSING

Data Wrangling Methodology

Cleaning and processing steps

- 01 Calculate missing-value percentages
- 02 Replace missing PayloadMass with the mean
- 03 Retain LandingPad nulls as meaningful no-pad cases
- 04 Count launch sites and orbit occurrences
- 05 Map unsuccessful outcomes to Class = 0
- 06 Map remaining outcomes to Class = 1

Result

90 rows
0 missing PayloadMass
87.8% success share

GitHub URL: <https://github.com/TristinWen/spacex-falcon9-landing-prediction>

Supporting visual from the completed analysis

EDA with Data Visualization: Methods and Chart Purpose

Evidence and interpretation

Scatter plots test relationships among flight number, payload, launch site, orbit, and landing outcome. A bar chart compares mean success by orbit type. A line chart measures the yearly success trend. Each result slide states the chart purpose and interprets the observed pattern.

Repository evidence

<https://github.com/TristinWen/spacex-falcon9-landing-prediction>

RUBRIC 1.8 - EDA VISUALIZATION SUMMARY

EDA with Data Visualization: Methods and Chart Purpose

Scatter plots	Test relationships among flight number, payload, site, orbit, and landing class
Bar chart	Compare average landing success across orbit types
Line chart	Show the yearly landing-success trend

Result slides follow in Rubric 1.11 order.

GitHub URL: <https://github.com/TristinWen/spacex-falcon9-landing-prediction>

Supporting visual from the completed analysis

EDA with SQL: Query Summary

Evidence and interpretation

SQL queries identify unique launch sites, CCA date-range records, total and average payload mass, the first successful ground-pad landing, booster success and failure totals, maximum-payload missions, monthly failure records, and landing-outcome rankings.

Repository evidence

<https://github.com/TristinWen/spacex-falcon9-landing-prediction>

RUBRIC 1.9 - SQL QUERY METHODS

EDA with SQL: Query Summary

- Unique launch sites and CCA site records
- Total and average payload mass
- First successful ground-pad landing date
- Boosters with successful drone-ship landings
- Mission success and failure totals
- Maximum-payload booster versions
- Monthly failure records and landing-outcome ranking

SQL result slides follow in Rubric 1.12 order.

GitHub URL: <https://github.com/TristinWen/spacex-falcon9-landing-prediction>

Supporting visual from the completed analysis

Interactive Visual Analytics: Folium and Plotly Dash

Evidence and interpretation


Folium maps launch-site coordinates, outcome markers, marker clusters, and distances to coastline, cities, highways, and railways. Plotly Dash filters records by site and payload range and dynamically updates a success pie chart and payload-versus-outcome scatter plot.

Repository evidence

<https://github.com/TristinWen/spacex-falcon9-landing-prediction>


RUBRIC 1.10 - INTERACTIVE VISUAL ANALYTICS SUMMARY

Interactive Visual Analytics: Folium and Plotly Dash



Folium

Launch-site markers, outcome clusters, and proximity distances.



Plotly Dash

Site dropdown, success pie charts, payload slider, and payload-versus-outcome scatter plots.

Folium: https://github.com/TristinWen/spacex-falcon9-landing-prediction/blob/main/notebooks/lab_jupyter_launch_site_location_completed.ipynb

Dash: <https://github.com/TristinWen/spacex-falcon9-landing-prediction/blob/main/dashboard/spacex-dash-app.py>

Supporting visual from the completed analysis


EDA Result: Flight Number vs. Launch Site

Evidence and interpretation

This scatter plot compares flight number and launch site, with point color representing landing outcome. Successful outcomes become more frequent in later flights across multiple sites, suggesting operational learning and maturation over time.

RUBRIC 1.11 - REQUIRED SCATTER PLOT

EDA Result: Flight Number vs. Launch Site



Purpose

Assess whether landing success changes with launch experience at each site.

Finding

Later flights show more successful outcomes across launch sites.

Supporting visual from the completed analysis


EDA Result: Payload Mass vs. Launch Site

Evidence and interpretation

This scatter plot compares payload mass across launch sites and distinguishes successful and failed landings. Launch sites operate over different payload ranges, so payload and site effects should be evaluated jointly rather than independently.

RUBRIC 1.11 - REQUIRED SCATTER PLOT

EDA Result: Payload Mass vs. Launch Site



Purpose

Compare payload ranges and landing outcomes by launch site.

Finding

Sites operate across different payload ranges and outcome patterns.

Supporting visual from the completed analysis


EDA Result: Success Rate by Orbit Type

Evidence and interpretation

The bar chart compares average Class by orbit type. ES-L1, GEO, HEO, and SSO show perfect success in this sample, while GTO has materially lower success. Orbit profile is therefore a useful predictive feature.

RUBRIC 1.11 - REQUIRED BAR CHART

EDA Result: Success Rate by Orbit Type



The bar chart compares mean Class by orbit.


ES-L1, GEO, HEO, and SSO show perfect success in this sample, while GTO is materially lower.

Supporting visual from the completed analysis

EDA Result: Flight Number vs. Orbit Type

Evidence and interpretation

This scatter plot compares flight sequence and orbit type. Several orbit categories show more successful outcomes at higher flight numbers, supporting the broader pattern of improvement as Falcon 9 operations matured.




EDA Result: Payload Mass vs. Orbit Type

Evidence and interpretation

This plot compares payload mass and landing outcome within each orbit category. The relationship varies by orbit and mission profile, indicating that payload effects interact with destination orbit rather than following one universal threshold.

RUBRIC 1.11 - REQUIRED SCATTER PLOT

EDA Result: Payload Mass vs. Orbit Type



Purpose

Compare mission payload ranges and outcomes across orbit types.

Finding

Payload-outcome relationships differ by orbit and mission profile.

Supporting visual from the completed analysis


EDA Result: Yearly Launch Success Trend

Evidence and interpretation

Average landing success rises from near zero in the early years to above 80% by 2019-2020, with a temporary decline in 2018. The overall trend provides strong evidence of technological and operational improvement.

RUBRIC 1.11 - REQUIRED YEARLY TREND

EDA Result: Yearly Launch Success Trend



Finding


Average landing success rises from zero in the early years to above 80% by 2019-2020, with a temporary dip in 2018.

Supporting visual from the completed analysis

EDA with SQL Results: Launch Sites and Payloads

Evidence and interpretation


SQL returned four unique launch sites and four CCA records in the requested date window. NASA CRS payload totals 28,596 kg; average F9 v1.1 payload is 4,256.4 kg; and twelve booster records share the maximum payload.



EDA with SQL Results: Success Rates and Rankings

Evidence and interpretation

SQL aggregates summarize mission outcomes and rank landing outcomes from 2010-2017. Successful drone-ship boosters in the 4,000-6,000 kg payload interval include B1022, B1026, B1021.2, and B1031.2.



EDA with SQL Results: Time Analysis

Evidence and interpretation

The first successful ground-pad landing occurred on 22 December 2015. Monthly queries also identify the 2015 drone-ship failures and their associated booster and launch-site records.

RUBRIC 1.12 - DATE, MONTH, AND TIME ANALYSIS

EDA with SQL Results: Time Analysis

First successful ground-pad landing

22 December 2015

2015 drone-ship failures

January - F9 v1.1 B1012 - CCAFS LC-40
April - F9 v1.1 B1015 - CCAFS LC-40

The ordered SQL analysis covers launch sites, payloads, success outcomes, rankings, dates, and months.

Supporting visual from the completed analysis


Folium Map Result: Launch Site Markers

Evidence and interpretation

The national Folium map shows all four launch-site labels at their coordinates. Marker clustering summarizes overlapping launch records and supports comparison of the geographic distribution of operations.

RUBRIC 1.13 - LAUNCH SITE MARKERS

Folium Map Result: Launch Site Markers



Map evidence

- Four launch-site labels
- Coastal locations in Florida and California
- Marker clusters summarize overlapping launch records

Supporting visual from the completed analysis


Folium Map Result: Launch Records and Proximity Analysis

Evidence and interpretation

Green and red markers distinguish successful and failed launch records. Distance lines show proximity to coastline, city, highway, and railway. The sites are close to the coast and transportation infrastructure, supporting range safety and logistics.

RUBRIC 1.13 - OUTCOME MARKERS AND PROXIMITY

Folium Map Result: Launch Records and Proximity Analysis



Completed evidence


- Green/red success and failure markers
- Distance to coastline
- Lines to city, railway, and highway points
- Launch sites are close to coastlines and transport infrastructure

Supporting visual from the completed analysis

Plotly Dash Result: Success Pie Chart

Evidence and interpretation


The all-sites Dash pie chart reports the contribution of each site to successful launches. KSC LC-39A has the largest share at 41.7%. The dropdown and payload slider demonstrate interactive filtering controls.



Plotly Dash Result: Payload vs. Launch Outcome Scatter Plot

Evidence and interpretation


Selecting CCAFS LC-40 updates both Dash charts. The site records show 26.9% successful launches and 73.1% failed launches, while the payload-outcome scatter plot responds to the selected site and payload interval.



Predictive Analysis: Multiple Models and Evaluation Results

Evidence and interpretation

GridSearchCV tuned logistic regression, support vector machine, decision tree, and K-nearest neighbors. Every classifier achieved 83.3% test accuracy. The decision tree achieved the strongest cross-validation score at 92.9%.




Predictive Analysis: Confusion Matrices

Evidence and interpretation


Four confusion matrices display true positives, true negatives, false positives, and false negatives for logistic regression, SVM, decision tree, and KNN. This comparison supplements accuracy by showing the type and distribution of classification errors.

RUBRIC 1.15 - CONFUSION MATRIX EVIDENCE


Predictive Analysis: Confusion Matrices




Logistic Regression



Decision Tree



Support Vector Machine



K-Nearest Neighbors

Supporting visual from the completed analysis

Best Model Explanation and Final Conclusion

Evidence and interpretation

The decision tree is selected because it combines the best cross-validation performance with interpretable nonlinear feature interactions. The model should support, not replace, engineering judgment. Future work should monitor drift, add weather and pad-condition variables, expand the data, and validate on newer launches. A key insight is that payload, orbit, site, and flight maturity interact rather than acting independently.

RUBRIC 1.15 - BEST MODEL, CONCLUSION, AND INSIGHTS

Best Model Explanation and Final Conclusion

Best cross-validation model

Decision Tree

92.9% validation accuracy

83.3% test accuracy

Why this choice

The decision tree captures nonlinear interactions among site, orbit, payload, booster block, reuse history, grid fins, and legs.

Creative insight

Use calibrated landing probabilities, not only binary labels, to estimate expected launch cost and bid risk. Expand the dataset and use repeated cross-validation before operational deployment.

Supporting visual from the completed analysis

Rubric Evidence Index in Required Order

Evidence and interpretation

The evidence index maps criteria 1.1-1.15 to the corresponding pages. It confirms the repository URL, completed PDF, executive summary, introduction, collection and wrangling, EDA and SQL methods and results, Folium, Dash, predictive evaluation, conclusion, creativity, and innovative insight.

FINAL QUALITY CHECK - RUBRIC 1.1 THROUGH 1.15

Rubric Evidence Index in Required Order

1.1	GitHub URL	Slide 2	1.9	EDA SQL summary	Slide 9
1.2	Completed presentation PDF	Slide 1-27	1.10	Interactive summary	Slide 10
1.3	Executive Summary	Slide 3	1.11	EDA result slides	Slide 11-16
1.4	Introduction	Slide 4	1.12	SQL result slides	Slide 17-19
1.5	SpaceX API	Slide 5	1.13	Folium result slides	Slide 20-21
1.6	Web Scraping	Slide 6	1.14	Plotly Dash result slides	Slide 22-23
1.7	Data Wrangling	Slide 7	1.15	Predictive Analysis	Slide 24-26
1.8	EDA Visualization summary	Slide 8			

Supporting visual from the completed analysis